



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Experiment No.4
Apply Stemming on the given Text input
Date of Performance:
Date of Submission:



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Aim: Apply Stemming on the given Text input.

Objective: Understand the working of stemming algorithms and apply stemming on the given input text.

Theory:

Stemming is a process of linguistic normalization, which reduces words to their word root word or chops off the derivational affixes. For example, connection, connected, connecting word reduce to a common word "conect".

Stemming is the process of producing morphological variants of a root/base word. Stemming programs are commonly referred to as stemming algorithms or stemmers. A stemming algorithm reduces the words “chocolates”, “chocolatey”, “choco” to the root word, “chocolate” and “retrieval”, “retrieved”, “retrieves” and reduces to the stem “retrieve”. Stemming is an important part of the pipelining process in Natural language processing. The input to the stemmer is tokenized words.

Applications of stemming :

1. Stemming is used in information retrieval systems like search engines.
2. It is used to determine domain vocabularies in domain analysis.

Porter's Stemmer Algorithm:

It is one of the most popular stemming methods proposed in 1980. It is based on the idea that the suffixes in the English language are made up of a combination of smaller and simpler suffixes. This stemmer is known for its speed and simplicity. The main applications of Porter Stemmer include data mining and Information retrieval. However, its applications are only limited to English words. Also, the group of stems is mapped on to the same stem and the output stem is not necessarily a meaningful word. The algorithms are fairly lengthy in nature and are known to be the oldest stemmer.

Example: EED -> EE means “if the word has at least one vowel and consonant plus EED ending, change the ending to EE” as ‘agreed’ becomes ‘agree’.



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Advantage: It produces the best output as compared to other stemmers and it has less error rate.

Limitation: Morphological variants produced are not always real words.

Code:

```
### Necessary Imports
from google.colab import drive
drive.mount('/content/drive')
Mounted at /content/drive
```

```
import nltk, re, pprint, string
from nltk import word_tokenize, sent_tokenize
string.punctuation = string.punctuation + '\"'+\"'+!'+-+'/'+'\\'+!'+-+'
string.punctuation = string.punctuation.replace('.', ' ')
file = open('/dataset.txt', encoding = 'utf8').read()
```

```
### Preprocess of the Data
file_nl_removed = ""
for line in file:
    line_nl_removed = line.replace("\n", " ")
    file_nl_removed += line_nl_removed
file_p = "".join([char for char in file_nl_removed if char not in
string.punctuation])
```

```
### Statistics of the Data
import nltk
nltk.download('punkt')
sents = nltk.sent_tokenize(file_p)
print("The number of sentences is", len(sents))

words = nltk.word_tokenize(file_p)
print("The number of tokens is", len(words))

average_tokens = round(len(words)/len(sents))
print("The average number of tokens per sentence is",
average_tokens)

unique_tokens = set(words)
print("The number of unique tokens are", len(unique_tokens))
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
```

CSDL7013: Natural Language Processing Lab



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

The number of sentences is 981

The number of tokens is 27361

The average number of tokens per sentence is 28

The number of unique tokens are 3039

```
### Building the N-Gram Model
```

```
import nltk
```

```
nltk.download('stopwords')
```

```
from nltk.util import ngrams
```

```
from nltk.corpus import stopwords
```

```
stop_words = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
```

```
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
unigram=[]
```

```
bigram=[]
```

```
trigram=[]
```

```
fourgram=[]
```

```
tokenized_text = []
```

```
for sentence in sents:
```

```
    sentence = sentence.lower()
```

```
    sequence = word_tokenize(sentence)
```

```
    for word in sequence:
```

```
        if (word == '.')
```

```
            sequence.remove(word)
```

```
        else:
```

```
            unigram.append(word)
```

```
    tokenized_text.append(sequence)
```

```
    bigram.extend(list(ngrams(sequence, 2)))
```

```
    trigram.extend(list(ngrams(sequence, 3)))
```

```
    fourgram.extend(list(ngrams(sequence, 4)))
```

```
#removes ngrams containing only stopwords
```

```
def removal(x):
```

```
    y = []
```

```
    for pair in x:
```

```
        count = 0
```

```
        for word in pair:
```

```
            if word in stop_words:
```

```
                count = count or 0
```

```
            else:
```

```
                count = count or 1
```

```
        if (count==1):
```

```
            y.append(pair)
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
return(y)

bigram = removal(bigram)
trigram = removal(trigram)
fourgram = removal(fourgram)
freq_bi = nltk.FreqDist(bigram)
freq_tri = nltk.FreqDist(trigram)
freq_four = nltk.FreqDist(fourgram)
print("Most common n-grams without stopwords removal and without add-1 smoothing: \n")
print ("Most common bigrams: ", freq_bi.most_common(5))
print ("\nMost common trigrams: ", freq_tri.most_common(5))
print ("\nMost common fourgrams: ", freq_four.most_common(5))
Most common n-grams without stopwords removal and without add-1 smoothing:

Most common bigrams: [(('said', 'the'), 209), (('said', 'alice'), 115), (('the', 'queen'), 65), (('the', 'king'), 60), (('a', 'little'), 59)]

Most common trigrams: [(('the', 'mock', 'turtle'), 51), (('the', 'march', 'hare'), 30), (('said', 'the', 'king'), 29), (('the', 'white', 'rabbit'), 21), (('said', 'the', 'hatter'), 21)]

Most common fourgrams: [(('said', 'the', 'mock', 'turtle'), 19), (('she', 'said', 'to', 'herself'), 16), (('a', 'minute', 'or', 'two'), 11), (('said', 'the', 'march', 'hare'), 8), (('will', 'you', 'wont', 'you'), 8)]

### Script for downloading the stopwords using NLTK
from nltk.corpus import stopwords
stop_words = set(stopwords.words('english'))

### Print 10 Unigrams and Bigrams after removing stopwords
print("Most common n-grams with stopwords removal and without add-1 smoothing: \n")
unigram_sw_removed = [p for p in unigram if p not in stop_words]
fdist = nltk.FreqDist(unigram_sw_removed)
print("Most common unigrams: ", fdist.most_common(10))
bigram_sw_removed = []
bigram_sw_removed.extend(list(ngrams(unigram_sw_removed, 2)))
fdist = nltk.FreqDist(bigram_sw_removed)
print("\nMost common bigrams: ", fdist.most_common(10))
Most common n-grams with stopwords removal and without add-1 smoothing:
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Most common unigrams: [('said', 462), ('alice', 385), ('little', 128), ('one', 101), ('like', 85), ('know', 85), ('would', 83), ('went', 83), ('could', 77), ('thought', 74)]

Most common bigrams: [(['said', 'alice'), 122), (['mock', 'turtle'], 54), (['march', 'hare'], 31), (['said', 'king'], 29), (['thought', 'alice'], 26), (['white', 'rabbit'], 22), (['said', 'hatter'], 22), (['said', 'mock'], 20), (['said', 'caterpillar'], 18), (['said', 'gryphon'], 18)]

```
### Add-1 smoothing
ngrams_all = {1:[], 2:[], 3:[], 4:[]}
for i in range(4):
    for each in tokenized_text:
        for j in ngrams(each, i+1):
            ngrams_all[i+1].append(j)
ngrams_voc = {1:set([]), 2:set([]), 3:set([]), 4:set([])}
for i in range(4):
    for gram in ngrams_all[i+1]:
        if gram not in ngrams_voc[i+1]:
            ngrams_voc[i+1].add(gram)
total_ngrams = {1:-1, 2:-1, 3:-1, 4:-1}
total_voc = {1:-1, 2:-1, 3:-1, 4:-1}
for i in range(4):
    total_ngrams[i+1] = len(ngrams_all[i+1])
    total_voc[i+1] = len(ngrams_voc[i+1])

ngrams_prob = {1:[], 2:[], 3:[], 4:[]}
for i in range(4):
    for ngram in ngrams_voc[i+1]:
        tlist = [ngram]
        tlist.append(ngrams_all[i+1].count(ngram))
        ngrams_prob[i+1].append(tlist)

for i in range(4):
    for ngram in ngrams_prob[i+1]:
        ngram[-1] = (ngram[-1]+1)/(total_ngrams[i+1]+total_voc[i+1])
### Prints top 10 unigram, bigram, trigram, fourgram after smoothing
print("Most common n-grams without stopword removal and with add-1 smoothing: \n")
for i in range(4):
    ngrams_prob[i+1] = sorted(ngrams_prob[i+1], key = lambda x:x[1], reverse = True)

print ("Most common unigrams: ", str(ngrams_prob[1][:10]))
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
print ("\nMost common bigrams: ", str(ngrams_prob[2][:10]))
print ("\nMost common trigrams: ", str(ngrams_prob[3][:10]))
print ("\nMost common fourgrams: ", str(ngrams_prob[4][:10]))
Most common n-grams without stopword removal and with add-1 smoothing:
```

```
Most common unigrams: [[('the',), 0.05598462224968249], [('and',),
0.02900490852298081], [('to',), 0.02478289225277177], [('a',),
0.02155631071293722], [('she',), 0.018467030515223287], [('it',),
0.018089451824391582], [('of',), 0.017471595784848797], [('said',),
0.015892630350461675], [('i',), 0.013764459547592077], [('alice',),
0.013249579514639755]]
```

```
Most common bigrams: [[('said', 'the'), 0.0053395713087035016],
[('of', 'the'), 0.0033308754354293268], [('said', 'alice'),
0.0029494774848076483], [('in', 'a'), 0.002491799944061634], [('and',
'the'), 0.002059548933357065], [('in', 'the'), 0.0020086958732741743],
[('it', 'was'), 0.0019069897531083933], [('to', 'the'),
0.0017798571029011671], [('the', 'queen'), 0.0016781509827353861],
[('as', 'she'), 0.0015764448625696051]]
```

```
Most common trigrams: [[('the', 'mock', 'turtle'),
0.001143837575064341], [('the', 'march', 'hare'),
0.0006819031697498955], [('said', 'the', 'king'),
0.0006599062933063505], [('said', 'the', 'hatter'),
0.00048393128175799036], [('the', 'white', 'rabbit'),
0.00048393128175799036], [('said', 'the', 'mock'),
0.0004399375288709003], [('said', 'to', 'herself'),
0.0004399375288709003], [('said', 'the', 'caterpillar'),
0.0004179406524273553], [('she', 'went', 'on'), 0.0003959437759838103],
[('she', 'said', 'to'), 0.0003959437759838103]]
```

```
Most common fourgrams: [[('said', 'the', 'mock', 'turtle'),
0.00043521782652217433], [('she', 'said', 'to', 'herself'),
0.0003699351525438482], [('a', 'minute', 'or', 'two'),
0.0002611306959133046], [('said', 'the', 'march', 'hare'),
0.00019584802193497845], [('will', 'you', 'wont', 'you'),
0.00019584802193497845], [('said', 'alice', 'in', 'a'),
0.00017408713060886974], [('as', 'well', 'as', 'she'),
0.00015232623928276102], [('in', 'a', 'great', 'hurry'),
0.00015232623928276102], [('in', 'a', 'tone', 'of'),
0.00015232623928276102], [('the', 'moral', 'of', 'that'),
0.00015232623928276102]]
```

Next word Prediction

```
str1 = 'after that alice said the'
str2 = 'alice felt so desperate that she was'
token_1 = word_tokenize(str1)
token_2 = word_tokenize(str2)
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
ngram_1 = {1:[], 2:[], 3:[]}    #to store the n-grams formed
ngram_2 = {1:[], 2:[], 3:[]}
for i in range(3):
    ngram_1[i+1] = list(ngrams(token_1, i+1))[-1]
    ngram_2[i+1] = list(ngrams(token_2, i+1))[-1]
print("String 1: ", ngram_1, "\nString 2: ", ngram_2)
String 1: {1: ('the',), 2: ('said', 'the'), 3: ('alice', 'said',
'the')}
String 2: {1: ('was',), 2: ('she', 'was'), 3: ('that', 'she', 'was')}
for i in range(4):

    ngrams_prob[i+1] = sorted(ngrams_prob[i+1], key = lambda x:x[1],
reverse = True)

pred_1 = {1:[], 2:[], 3:[]}
for i in range(3):
    count = 0
    for each in ngrams_prob[i+2]:
        if each[0][-1] == ngram_1[i+1]:
            count +=1
            pred_1[i+1].append(each[0][-1])
            if count ==5:
                break
    if count<5:
        while(count!=5):
            pred_1[i+1].append("NOT FOUND")
#if no word prediction is found, replace with NOT FOUND
    count +=1
for i in range(4):
    ngrams_prob[i+1] = sorted(ngrams_prob[i+1], key = lambda x:x[1],
reverse = True)

pred_2 = {1:[], 2:[], 3:[]}
for i in range(3):
    count = 0
    for each in ngrams_prob[i+2]:
        if each[0][-1] == ngram_2[i+1]:
            count +=1
            pred_2[i+1].append(each[0][-1])
            if count ==5:
                break
    if count<5:
        while(count!=5):
            pred_2[i+1].append("\0")
```




Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
count +=1
print("Next word predictions for the strings using the probability
models of bigrams, trigrams, and fourgrams\n")
print("String 1 - after that alice said the-\n")
print("Bigram model predictions: {}\nTrigram model predictions:
{}\nFourgram model predictions: {}" .format(pred_1[1], pred_1[2],
pred_1[3]))
print("String 2 - alice felt so desperate that she was-\n")
print("Bigram model predictions: {}\nTrigram model predictions:
{}\nFourgram model predictions: {}" .format(pred_2[1], pred_2[2],
pred_2[3]))
Next word predictions for the strings using the probability models of
bigrams, trigrams, and fourgrams
```

String 1 - after that alice said the-

```
Bigram model predictions: ['queen', 'king', 'mock', 'gryphon',
'hatter']
Trigram model predictions: ['king', 'hatter', 'mock', 'caterpillar',
'gryphon']
Fourgram model predictions: ['NOT FOUND', 'NOT FOUND', 'NOT FOUND',
'NOT FOUND', 'NOT FOUND']
```

String 2 - alice felt so desperate that she was-

```
Bigram model predictions: ['a', 'the', 'not', 'going', 'that']
Trigram model predictions: ['now', 'quite', 'a', 'beginning',
'walking']
Fourgram model predictions: ['now', 'walking', 'quite', 'ready',
'losing']
```

Conclusion:

Stemming is a text normalization process that reduces words to their root or base form. It helps in handling variations of words. For English text, stemming works well in removing suffixes but can lead to errors. In Indian languages, stemming can be more complex due to diverse word structures and scripts. Conclusion: Stemming is valuable in English text analysis but requires language-specific algorithms for accurate results in Indian languages due to their linguistic complexity